


COMMUNICATION NETWORKS

DR LINA BARAKAT

School of Computer Science and Electronic Engineering

University of Essex

Email: lina.barakat@essex.ac.uk

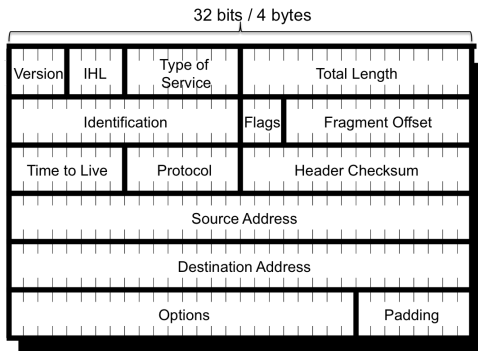


- Also called the internet layer or the network layer
- Provides a *virtual network* view (hides the underlying physical network)
- **Internet Protocol (IP)** is the most important internetwork-layer protocol
- Other internetwork-layer protocols:
 - Internet Control Message Protocol (ICMP)
 - Internet Group Management Protocol (IGMP)
 - Address Resolution Protocol (ARP)
 - Dynamic Host Configuration Protocol (DHCP)

- IP sends messages across the internet: provides a *routing function* that attempts to deliver messages to the right destination
- **IP datagram** is the basic unit of information transmitted in an IP network
 - messages may be divided into fragments, which are sent individually
 - IP adds a header to each fragment to form an IP datagram
- IP is an unreliable protocol: fragments might be lost, arrive out of order, or be duplicated
- Higher layer protocols add reliability to IP


IP Datagram Header

- IP Datagram = IP header + data (from higher-layer protocols)
- An IP header has a minimum length of 20 bytes



IP Datagram Header

- **Version field.** The IP protocol version being used
 - 4 bits
 - e.g. IPv4 is 0100
 - e.g. IPv6 is 0110
- **Internet Header Length (IHL) field.** The length of this IP header in 32-bit words
 - 4 bits
 - e.g. IHL value 6 (0110) → header is $6 * 32 = 192$ bits
 - Minimum IHL value is 5 (0101)
- **Service Type field.** Specifies how the datagram should be handled in transmission (the quality of service requested)
 - 8 bits



IP Datagram Header

- **Total length field.** Total number of bytes/octetets (multiples of 8 bits) that the IP datagram (header plus data) takes up
 - 16 bits
 - maximum length: 65,535 octets, minimum length: 20 octets
- **Identification field.** A unique number assigned to a datagram fragment to help reconstructing message from fragments
 - 16 bits
 - Every fragment of the same message has the same identification number
- **Flag fields.** Informational bits
 - 3 bits for 3 flags: Reserved, Don't Fragment (DF), and More Fragments (MF)
 - Reserved bit must be zero
 - DF indicates if fragmentation of this data is allowed (0), or not allowed (1)
 - MF indicates if this is the last fragment of the datagram (0), or there are still more fragments to follow (1)

IP Datagram Header

- **Fragment offset field.** Specifies where in the original message the fragment in this datagram starts
 - 13 bits
 - It is the position of this fragment from the start of the original message counted in 8 bytes units
- **Time to live field (8 bits).** How long this datagram is allowed to remain in the system before being deleted
- **Protocol field (8 bits).** The host-to-host transport layer protocol being used by the message sent over IP, e.g. TCP: 6
- **Header checksum field (16 bits).** Used to ensure integrity of the header in transmission

IP Datagram Header

- **Source field (32 bits).** IP address of sending/source host, e.g.
137.73.9.232
- **Destination field (32 bits).** IP address of destination host (intended recipient)
- **Options field.** Optional arguments usable by IP processing software
- **Padding field.** 0s to make header up to multiple of 32 bits

Fragmentation

- Each network has a limit on the size of message it can carry: **the Maximum Transmission Unit (MTU)** of the network
- The sender fragments messages to fit the local MTU
- In IPv4, the fragments may be further fragmented on moving from one network to another with smaller MTU
- Fragments are reassembled at the destination

Fragmentation Algorithm

- ➊ The DF flag bit is checked to see if fragmentation is allowed
 - If DF is set, but the message needs to be fragmented to move onto a network with a lower MTU, an error message is sent by ICMP
 - ICMP is part of IP and used to report errors and other info regarding IP processing back to the datagram's source
 - ICMP code 4: fragmentation needed and DF set
- ➋ Based on the MTU value, the data field is split into two or more parts
 - every newly created data portion must be a multiple of 8 bytes (64 bits) except the last one
- ➌ Each data portion is placed in an IP datagram.
- ➍ Each of the fragmented datagrams is forwarded as a normal IP datagram
 - The fragments can traverse different routers to the intended destination

Fragmentation Algorithm

The **headers** of fragmented datagrams are minor modifications of the original datagram's header:

- The MF flag bit is set in all fragments except the last
- The fragment offset field in each is set to the location this data portion occupied in the original datagram, relative to the beginning of the original datagram
- If options were included in the original datagram, some of the options may be copied to each of the fragments depending on the Copied flag in each option field
- The header length field of the fragment is set
- The total length field of the fragment is set
- The header checksum field is re-calculated

- Transmission Control Protocol (TCP)
- Key Features:
 - **Connection oriented**: implements mechanisms to setup and tear down a full duplex connection between end points
 - **Reliable**: implements mechanisms to guarantee error free and ordered delivery of information
 - **Flow and Congestion controlled**: implements mechanisms to control traffic

Ports and Sockets

- A process identifies itself to the TCP protocol by one or more **ports**
- A port is a 16-bit number used to identify to which application/process the message should be delivered
- Some ports are reserved for specific applications
 - e.g. FTP: 20/21, HTTP: 80, SMTP (e-mail): 25
- A **socket** is the combination of a host's IP address and a port number
 - E.g. 137.73.9.232 : 8080
 - Every communication in TCP is between two sockets, i.e. two hosts using particular ports

TCP Connection Set-up

Connection is set up via a **handshake** that involves three steps:

- Sender sends a synchronise message (SYN) to the receiver
- Receiver sends a message back acknowledging the synchronise (SYN ACK), and giving permission for communication to take place
- Sender sends a message acknowledging the acknowledgement (ACK)


TCP Connection tear-down

Connection is closed via a **handshake** that involves three steps:

- Sender sends a finalise message (FIN) to the receiver
- Receiver responds with an acknowledgement of the finalise (FIN + ACK)
- Finally, the sender responds with an acknowledgement of the acknowledgement (ACK)

TCP Concepts for Reliability and Flow Control


Problem with Unreliable Protocols: Lack of Feedback



TCP Concepts for Reliability and Flow Control

Basic Reliability: Positive Acknowledgment with Retransmission (PAR)

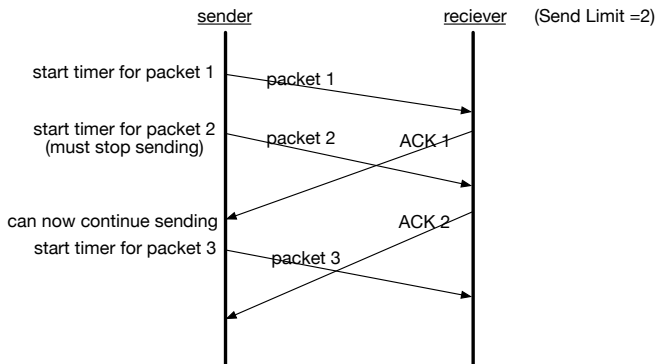
- Drawback: sender cannot send a second message until the first has been acknowledged



TCP Concepts for Reliability and Flow Control

Improved PAR: with Message Identification and Send Limits

- multiple messages can be sent without waiting for acknowledgment
- ensures better use of the network bandwidth (**efficiency**)
- *send limit* specifies maximum number of unacknowledged messages allowed from sender at one time (**basic flow control**)



TCP Concepts for Reliability and Flow Control


TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments

- TCP is a stream oriented protocol
- Every byte in a message sent from a sender to a receiver has a **sequence number**
 - Sender communicates its **initial sequence number (ISN)** in synchronisation
 - Sequence number for 1st byte of message is $ISN+1$, for 2nd byte is $ISN+2$, ...
- TCP divides continuous byte stream into **segments**
 - a TCP segment only carries the sequence number of the first byte in the segment

TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments


- **Reliability.** The receiver sends an **acknowledgement** back to the sender for every segment it receives
 - An acknowledgement states that the receiver has received all data in the message before a given sequence number
 - A segment is retransmitted on **time out**, or on acknowledgment **repetition**



TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments


- **Reliability.** The receiver sends an **acknowledgement** back to the sender for every segment it receives
 - An acknowledgement states that the receiver has received all data in the message before a given sequence number
 - A segment is retransmitted on **time out**, or on acknowledgment **repetition**



TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments


- **Reliability**. The receiver sends an **acknowledgement** back to the sender for every segment it receives
 - An acknowledgement states that the receiver has received all data in the message before a given sequence number
 - A segment is retransmitted on **time out**, or on acknowledgment **repetition**



TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments


- **Reliability**. The receiver sends an **acknowledgement** back to the sender for every segment it receives
 - An acknowledgement states that the receiver has received all data in the message before a given sequence number
 - A segment is retransmitted on **time out**, or on acknowledgment **repetition**



TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments


- **Reliability.** The receiver sends an **acknowledgement** back to the sender for every segment it receives
 - An acknowledgement states that the receiver has received all data in the message before a given sequence number
 - A segment is retransmitted on **time out**, or on acknowledgment **repetition**



TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments

- **Reliability**. The receiver sends an **acknowledgement** back to the sender for every segment it receives
 - An acknowledgement states that the receiver has received all data in the message before a given sequence number
 - A segment is retransmitted on **time out**, or on acknowledgment **repetition**



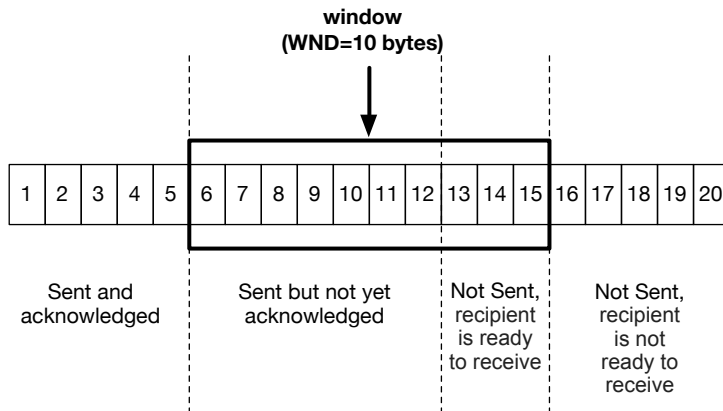
TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments

- **Flow Control.** The amount of data that the receiver can receive is called the **window size (WND)**
 - It is the maximum number of unacknowledged bytes allowed from sender at one time
 - Determined initially by the receiver when the connection is established, but can vary during data transfer
 - Each ACK message will include the window size that the receiver is ready to deal with at that particular time
 - **Usable Window:** amount of the window that the sender is still allowed to send at any point in time

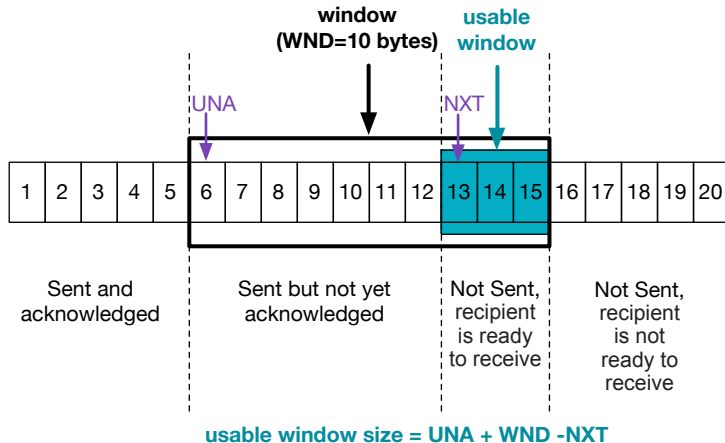
TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments



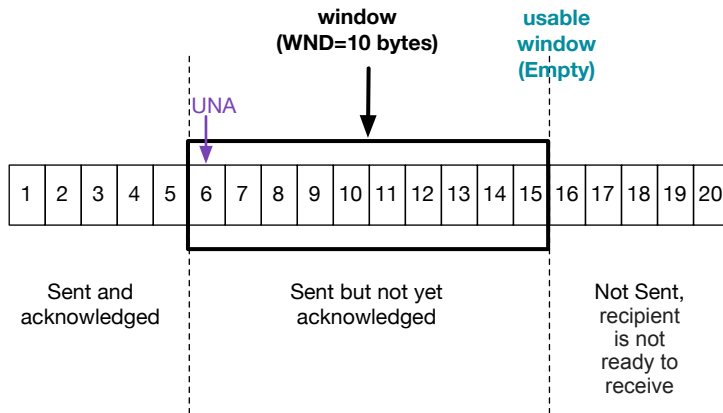
TCP Concepts for Reliability and Flow Control

TCP's Sliding Window System: Similar Conceptually to Improved PAR, but with Some Adjustments



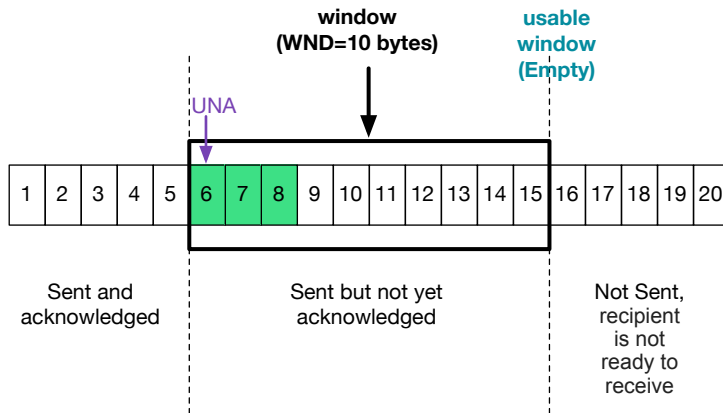
TCP's Sliding Window System: Example

Sending bytes in the usable window:



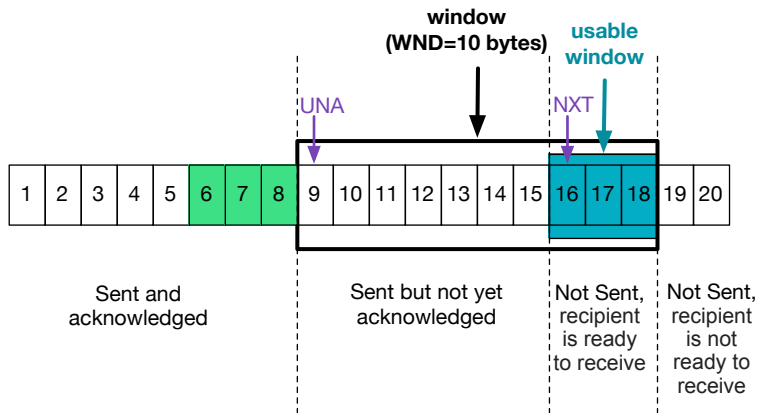
TCP's Sliding Window System: Example

Processing acknowledgments (ACK=9)

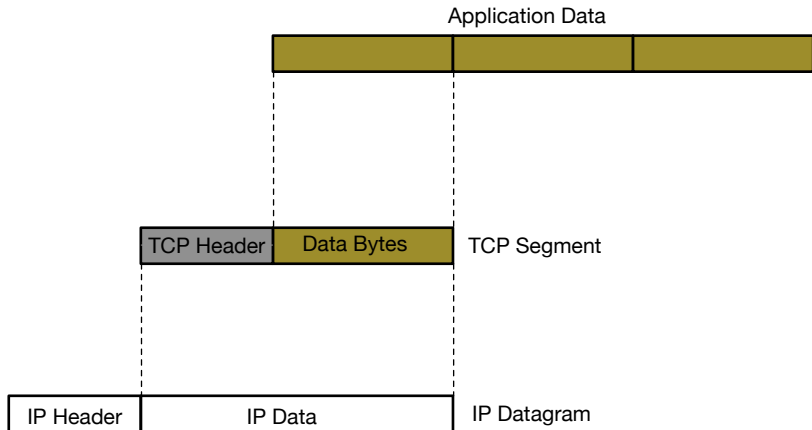


TCP's Sliding Window System: Example

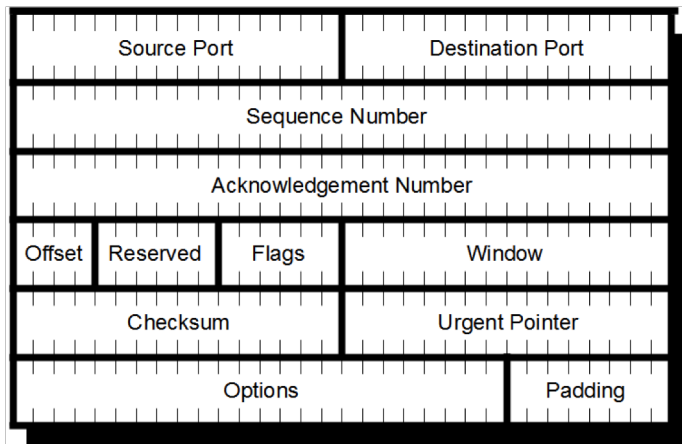
Processing acknowledgments (ACK=9)



TCP and IP



TCP Header



- Source Port: the port of the segment sender
- Destination Port: the port of the receiver
- Sequence Number: the sequence number of the start of the segment, or the Initial Sequence Number (if a synchronise message)
- Acknowledgement Number: the sequence number before which the receiver has all the message data
- Data Offset: The length of the TCP header in 32-bit words
- Reserved: Not currently used (all 0s)

- URG Flag: Marks that this message contains urgent data
- ACK Flag: Marks that this is an acknowledgement
- PSH Flag: Marks that this data was pushed
- RST Flag: Marks that this is a reset message
- SYN Flag: Marks that this is a synchronise message or acknowledgement of a synchronise
- FIN Flag: Marks that this is a finalise message or acknowledgement of a finalise

- Window: The current acceptable window size (sent in acknowledgement messages)
- Checksum: A checksum over the segment, used to check for corruption
- Urgent Pointer: Position of where the urgent data ends inside the segment
- Options: Various TCP options (varies in length)

Useful Resources

- TCP (RFC 793)
- TCP and IP (RFC 879)
- <http://www.tcpipguide.com>